

# High Performance Machine Learning

## Lab 3

Ankit Rajvanshi  
ar7996

### Chatbot Seq-2-Seq Model

Wandb Project: <https://wandb.ai/ar7996/hpml-lab-3?nw=nwuserar7996>

#### Hyperparameter Sweeps:

Strategy: Random Search

Total Runs: 25

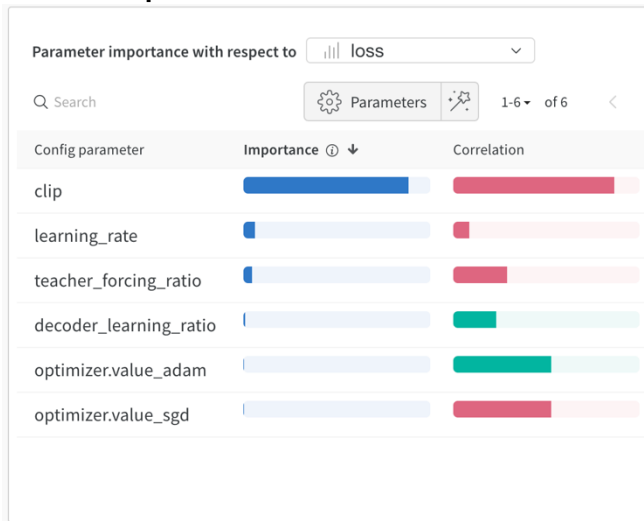
<input type="checkbox"/>	Name (25 visualized)	ID	Notes	Use	Tag	Cre	Run	Sweep	clip	decode	learn	optimi	teache	iteratic	loss
	sleek-sweep-21	wmnkkrm	Add notes	ar7996		1h ago	3m 58s	iwg43qvh	100	10	0.001	sgd	0	4000	NaN
	lucky-sweep-20	odvx6vak	Add notes	ar7996		1h ago	3m 59s	iwg43qvh	25	1	0.0005	sgd	0	4000	4.272
	electric-sweep-19	0hl3y43r	Add notes	ar7996		1h ago	4m 1s	iwg43qvh	100	10	0.001	sgd	0	4000	NaN
	fiery-sweep-18	1r6sz3jz	Add notes	ar7996		1h ago	3m 59s	iwg43qvh	50	3	0.00025	sgd	0	4000	3.744
	rich-sweep-17	fdk0wnzj	Add notes	ar7996		1h ago	3m 42s	iwg43qvh	50	3	0.0005	sgd	0.5	4000	3.929
	playful-sweep-16	aka90sa9	Add notes	ar7996		1h ago	3m 24s	iwg43qvh	50	5	0.00025	adam	1	4000	2.032
	ancient-sweep-15	7yxhjld	Add notes	ar7996		1h ago	4m	iwg43qvh	100	3	0.00025	adam	0	4000	3.77
	smooth-sweep-14	woqfwfm	Add notes	ar7996		1h ago	3m 42s	iwg43qvh	0	10	0.0001	adam	0.5	4000	8.975
	gallant-sweep-13	24u21q5a	Add notes	ar7996		1h ago	3m 24s	iwg43qvh	50	3	0.001	adam	1	4000	8.976
	blooming-swee...	61jalbib	Add notes	ar7996		1h ago	3m 48s	iwg43qvh	0	3	0.00025	adam	0.5	4000	8.976
	exalted-sweep-11	1xbmmrlk	Add notes	ar7996		1h ago	3m 26s	iwg43qvh	0	10	0.00025	sgd	1	4000	8.969

From above results, we can observe that run with id **aka90sa9** has least loss of **2.032**.

Hyperparameters value of above model:

1. Optimizer: ADAM
2. Learning Rate: 0.00025
3. Clip: 50
4. Teacher Forcing Ratio: 1
5. Decoder Learning Ratio: 5


## Feature Importance:



From above panel, we can observe that value of **clip** has very high effect on the **loss** i.e if value of clip is high, loss will be less (Value of correlation is negative).


Moreover, we can see that both “ADAM” and “SGD” has high correlation, but ADAM has positive correlation and hence ADAM will be better choice.

The importance metric is very high for clip with respect to other hyperparameters which shows that if clip value is very less, no matter what the value of other parameters is, loss will be greater. It is evident by following diagram.



Variation of loss with iterations

For each run, there were 4000 iterations.



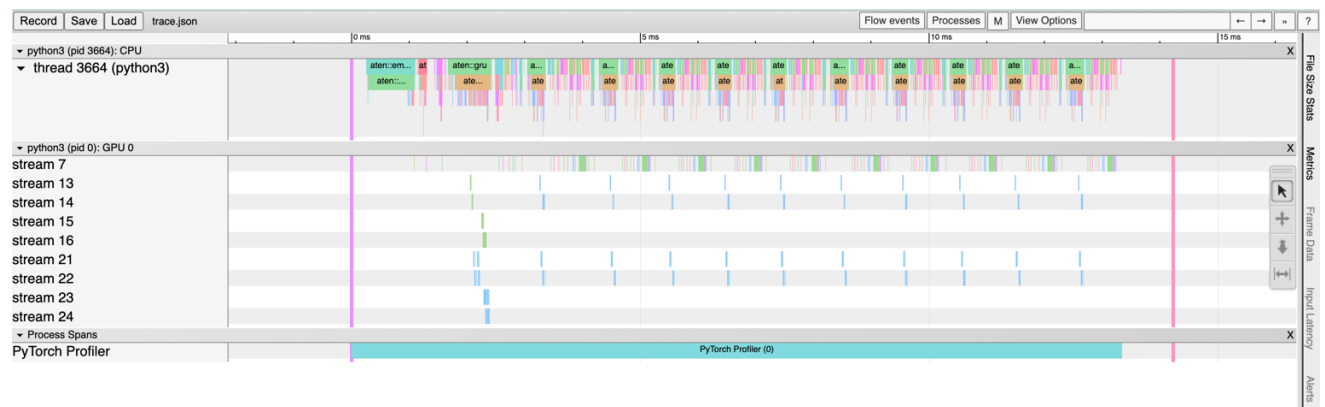
Measurement of time and memory Consumption of model’s operators

Full measurement is present in *profiler.txt* file.

		Name	Self CPU	CPU total	CPU time avg	CUDA total	CUDA time avg	CPU Mem	Self CPU Mem	CUDA Mem	Self CUDA Mem
ous namespace)		aten::empty	523.000us	523.000us	6.226us	0.000us	0.000us	24 b	24 b	335.39 Mb	335.39 Mb
		aten::embedding	4.636ms	6.319ms	574.455us	56.000us	5.091us	0 b	0 b	24.00 Kb	0 b
		aten::reshape	27.000us	41.000us	3.727us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::view	31.000us	31.000us	1.292us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::index_select	1.028ms	1.631ms	148.273us	56.000us	5.091us	0 b	0 b	24.00 Kb	0 b
		aten::resize_	99.000us	99.000us	9.000us	0.000us	0.000us	0 b	0 b	24.00 Kb	24.00 Kb
		cudaLaunchKernel	33.995ms	33.995ms	157.384us	209.000us	0.968us	0 b	0 b	0 b	0 b
		aten::to	0.000us	0.000us	0.000us	56.000us	5.091us	0 b	0 b	0 b	0 b
		aten::_pack_padded_sequence	3.000us	3.000us	3.000us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::_pack_padded_sequence	51.000us	964.000us	964.000us	3.000us	3.000us	16 b	0 b	4.00 Kb	0 b
		aten::slice	201.000us	210.000us	16.154us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::as_strided	69.000us	69.000us	0.476us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::cat	1.278ms	1.735ms	55.968us	133.000us	4.290us	0 b	0 b	54.00 Kb	54.00 Kb
		aten::narrow	7.000us	16.000us	16.000us	0.000us	0.000us	0 b	0 b	0 b	0 b
		cudaMemcpyAsync	49.000us	49.000us	24.500us	0.000us	0.000us	0 b	0 b	0 b	0 b
		Memcpy DtoD (Device -> Device)	0.000us	0.000us	0.000us	6.000us	3.000us	0 b	0 b	0 b	0 b
		aten::select	101.000us	113.000us	5.136us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::item	8.000us	10.000us	5.000us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::_local_scalar_dense	4.000us	4.000us	2.000us	0.000us	0.000us	0 b	0 b	0 b	0 b
		aten::zeros	27.000us	1.613ms	537.667us	2.000us	0.667us	0 b	0 b	8.00 Kb	0 b

## PyTorch Profiler Tracing

Trace file has been saved as *trace.json*



## TorchScript Seq-2-Seq Model

**Q1. Explain the differences between tracing and scripting and how they are used in TorchScript?**

**Answer:**

**Tracing:** In case of tracing, function takes model and sample data as inputs and records the computations and builds graph-based function. Hence, only computations which were occurred for given data will be recorded. It is incapable of capturing data-driven control flow.

**Scripting:** In case of scripting, functions require only model without sample data. It converts model code into TorchScript (which is a subset of python language), including all control flows.

For a simple model without any control flow in its logic, we can easily use `torch.jit.trace()` without any modification to the model to trace the model and it will be converted to TorchScript. Whereas in case of *scripting*, we may need to rewrite the code to make sure it fits the TorchScript syntax. Once the code is rewritten adhere to TorchScript requirements, we can use `torch.jit.script()` to convert to TorchScript.

In case of *trace*, we need to set the device and dropout layers of model to test mode before tracing the model because traced version of model does not have functionality to perform these two operations. But in case of *scripting*, we can set the device and dropout layers to test mode just before the inferencing as we do in case of normal eager mode.

**Q2. Explain the changes needed in the chatbot model to allow for scripting.**

**Answer:**

In the given model, there are three sub modules: *Encoder*, *Decoder* and *GreedySearchDecoder*. Third sub module *GreedySearchDecoder* requires scripting as it has input based control flow in its logic.

Changes required in *GreedySearchDecoder*:

1. **Add additional arguments in its constructor for *decoder\_n\_layers*.** Earlier, it was using fetching this value from decoder but since we are using traced version of decoder we will not be able to access that value anymore and hence we need to pass this to its constructor for it to use.

```
#Modified GreedySearchDecoder for scripting the module

class GreedySearchDecoderScript(torch.jit.ScriptModule):
    def __init__(self, encoder, decoder, decoder_n_layers):
```

2. **Need to add more attributes.** Earlier, we were accessing various variable available in global scope of our python environment, but TorchScript version do not have access to those variables, and we need to save value of those variables from global scope to attributes of the model class. (*\_SOS\_token*, *\_device*, *\_decoder\_n\_layers*)

```
self._device = device
self._SOS_token = SOS_token
self._decoder_n_layers = decoder_n_layers
```

3. **Store above attributes as constants.** We need to add above attributes to special list called *\_\_constants\_\_* so that they can be used as literal values when constructing the graph in the forward method.

```
__constants__ = ['_device', '_SOS_token', '_decoder_n_layers']
```

4. **Enforce types of *forward* method arguments.** By default, TorchScript assume all parameters of function as tensor. Hence, in case we need to pass any argument with different type like int in our case, we need to specify the type in python function.

```
@torch.jit.script_method
def forward(self, input_seq : torch.Tensor, input_length : torch.Tensor, max_length : int):
```

### Q3: Compare Latency

Initial Run:

	Latency on CPU (ms)	Latency on GPU (ms)
Pytorch	37.62	10.56
Torchscript	27.4	14.84
SpeedUp	1.37	0.71

Second Run:

	Latency on CPU (ms)	Latency on GPU (ms)
Pytorch	37.62	10.56
Torchscript	15.81	7.9
SpeedUp	2.38	1.34